

# Reaching new heights with GeoFM

---

Team: Attention\_Plzzz

Dingqi Ye, Daniel Kiv, Wen Zhou, Wei Hu, Ayush Khot

July 2, 2026

CyberGIS Center for Advanced Digital and Spatial Studies  
University of Illinois at Urbana-Champaign

# Roadmap

Task & Inputs

Data observations

Architecture

Design 1 — Dense pixel backbone

Design 2 — Coarse token conditioning

Design 3 — Split-trunk head

Design 4 — Two-stage training

Loss

Inference & Results

Final Submission

Summary

## Task & Inputs

---

# The task

Per  $256 \times 256$  tile, predict 4 channels:

[building, vegetation, water, height (m)]

ch 0–2 = binary presence    ch 3 = continuous height.

|        | IoU bld | IoU tree | IoU water | RMSE bld <sub>H</sub> | RMSE veg <sub>H</sub> |
|--------|---------|----------|-----------|-----------------------|-----------------------|
| weight | 25%     | 15%      | 15%       | 25%                   | 20%                   |

Score =  $\sum_i w_i \text{IoU}_i + \sum_j w_j \max(0, 1 - \text{RMSE}_j / X_j)$ ,  $X_{\text{bld}}=3\text{m}$ ,  $X_{\text{veg}}=5\text{m}$ .

Everyone gets the *same frozen embeddings*. No raw imagery  $\Rightarrow$  a score gap is a **method** gap.

## The inputs: two families of embedding

| Embedding       | Source        | Shape              | Note                |
|-----------------|---------------|--------------------|---------------------|
| AlphaEarth      | multi-sensor  | $64 \times 256^2$  | dense, full-res     |
| Tesseract       | multi-sensor  | $128 \times 256^2$ | dense, full-res     |
| TerraMind S1/S2 | SAR / optical | $768 \times 16^2$  | post-encoder tokens |
| Thor S1/S2      | SAR / optical | $768 \times 16^2$  | raw scale $\pm 17k$ |

### Dense, full-res

per-pixel features  $\rightarrow$  boundaries & texture (drives IoU).

### Coarse tokens ( $16^2$ )

$16 \times$  downsampled  $\rightarrow$  global semantics, no detail.

This split — dense vs. coarse — dictates the whole design.

## The labels (and how the metric reads them)

Target tensor per pixel:  $[c_{\text{bld}}, c_{\text{veg}}, c_{\text{water}}, h/30]$

- ch 0–2: **coverage fractions**  $c \in [0, 1]$  (a 10 m pixel is rarely pure).
- ch 3: nDSM **height** in metres, normalised by 30.
- `valid_mask`: separate global (nodata) + height (nDSM holes) masks.

How the metric binarises them — and why it matters later:

- **IoU** compares  $\text{pred} > 0.5$  vs. official  $\text{GT} = \text{coverage} > 0.10$ .
- **RMSE** is masked *per class* (only pixels of that class count).

## Data observations

---

## Two label defects we found — and masked

Auditing the GT surfaced two *systematic* defects. Both are handled by the two-channel `valid_mask`, so a defect never becomes a wrong training signal.


| Defect                    | What we see                                                  | How we handle it                                      |
|---------------------------|--------------------------------------------------------------|-------------------------------------------------------|
| <b>nDSM holes</b>         | height = 0 on labelled objects<br>(often whole sub-tiles)    | drop px from <i>height</i> loss;<br>keep presence     |
| <b>Deleted footprints</b> | building label empty where<br>nDSM + embeddings say building | drop px from <i>presence/seg</i> loss;<br>keep height |

- `valid_mask` = 2 channels: (0) global/presence, (1) height.
- Each defect zeroes *only* its channel  $\Rightarrow$  the other task keeps learning.



# Observation A: nDSM holes


height label



Top: nDSM height as stored (the label writes a literal 0 → dark). Bottom: red = labelled land cover with height = 0 (up to ~40% of px). **Fits:** `ndsm_hole = height==0 & has_landcover` → dropped from height RMSE only; presence kept.

## Observation B: deleted building footprints

presence label




Footprint empty (top) where the nDSM (mid) shows buildings; a density-gap detector ( $\text{blob} \geq 50\text{px}$ ) recovers the region (red, bottom). **Fits:** dropped from presence/seg loss, height kept — never punish a *correct* building.

# Architecture

---

# Architecture at a glance




**Two-stage training.** Stage 1 (coupled, 80 ep) trains everything  $\Rightarrow$  a *shared init*. Stage 2 = **dual purify** off it: **seg-purify** (freeze height-grad)  $\Rightarrow$  sharper ch 0-2 and **height-purify** (freeze seg-grad)  $\Rightarrow$  ch 3; a +dDice continuation of seg-purify adds a 2nd, decorrelated seg source.

■ dense pixel ■ coarse token ■ split head ■ output / two-stage

## Design 1 — Dense pixel backbone

---



- **Fits:** both are dense  $256^2 \rightarrow$  a U-Net++ yields sharp, IoU-friendly masks.
- **Fits:** they differ in reliability  $\rightarrow$  the gate trusts each per pixel; bias init +4  $\Rightarrow$  starts AlphaEarth-dominant.

ViT / SegFormer / HRNet / UPerNet backbones were all tried and lost.

## U-Net++ (base\_ch=48, LightUNet base)

- 4 levels:  $48 \rightarrow 96 \rightarrow 192 \rightarrow 384$ .
- $3 \times$  MaxPool down ( $256 \rightarrow 32$ ), bilinear up + *nested* dense skips.
- DoubleConv blocks, BatchNorm.
- `forward_features`  $\rightarrow$  48-ch map at full res.


## Tessera entry & gate

- ChannelCalibration: per-channel affine  $x \cdot s + b$  (re-normalise the GFM's arbitrary scales).
- Gate = *rich*  
conv  $\rightarrow$  GN  $\rightarrow$  GELU  $\rightarrow$  conv, zero-init weights, bias +4 (tied:  $g$  and  $1-g$ ).

## Design 2 — Coarse token conditioning

---





- Each source projects to a shared context, cross-attends, then adds a zero-init residual (FiLM + additive + gate).
- **Fits:** coarse tokens can't make pixels  $\rightarrow$  they *modulate* the dense feature, never replace it.
- **Fits:** SAR vs. optical see different physics  $\rightarrow$  attention lets them inform each other.

# Token fusion internals

## Attend ( $16 \times 16$ )

- $1 \times 1$  proj:  $768 \rightarrow$  ctx 96 per source.
- 4-head self-attention over the  $4 \times 256$  tokens.
- learned modality embedding + 2D positional MLP.
- **zero-init** output proj  $\Rightarrow$  identity at start.

## Inject ( $256 \times 256$ )

- bilinear up  $16^2 \rightarrow 256^2$ , then FiLM+add+gate.
- all fusion convs **zero-init**.
- run in fp32 for stability.
- clamps: token/ctx  $\pm 50$ , FiLM/add  $\pm 4$ .


## Respecting each embedding's nature

- **Selective calibration.** Calibrate TerraMind (`indices=[0,1]`); leave Thor *uncalibrated + clamped*.  
**Fits:** Thor spans  $\pm 17k$  — calibrating it collapses water IoU.
- **Zero-init injection.** At init  $F_{\text{out}} \equiv F_{\text{pixel}}$ .  
**Fits:** boots as a pixel-only net, then *earns* token influence; a weak source degrades gracefully.

Attribution: dropping any of the 4 sources, or any of the 3 pathways, hurt the score.

## Design 3 — Split-trunk head

---



- Two trunks off  $F_{out}$ , sharing *no* head weights.
- **Fits:** IoU wants crisp masks, RMSE wants smooth regression — a shared trunk makes their gradients fight; splitting lets each shape its features.
- Presence head *is* the metric (BCE on coverage @ 0.5), not a reparam of fractions.

# Head internals & the gradient cut


## Trunks & heads

- `trunk = ConvGNAct + residual + GELU`.
- presence `split_all`: independent head per class (`branch_ch` 48, tower depth 2).
- height: 3 independent branches (`hidden` 96, depth 2).

## Soft gradient decouple

- forward identity:  
 $s t + (1-s) t.\text{detach}()$ .
- $s = 1$  coupled,  $s = 0$  hard cut.
- this is the knob the two-stage recipe flips.

## Presence-gated height specialists




$$H = p_{\text{fg}} (w_b h_{\text{bld}} + w_v h_{\text{veg}}) + (1 - p_{\text{fg}}) h_{\text{base}}, \quad p_{\text{fg}} = 1 - (1 - p_b)(1 - p_v)$$

- **Fits:** the metric masks RMSE *per class* — each specialist is L1-trained, and trusted, only on its own class's pixels.
- Gate uses *alpha* presence logits; each  $h$  comes from the soft-bin head (next slide).

Each specialist emits **64 bin logits**, not a scalar; the height is the **soft-argmax** (expectation over bin centers):

$$\hat{h} = \sum_k \text{softmax}(\ell)_k c_k$$



- **One shared grid:** log-spaced  $[0, 80]$  m, 64 bins — base/bld/veg share the centers (separate conv projections).
- bin-CE makes the distribution *commit* to a bin rather than collapse to a safe mean (vs. L1-only).
- Non-negative by construction ( $c_k \geq 0$ ).


**Refuted:** per-class grids (bld  $\log[0, 160]$ , veg  $linear[0, 40]$ ) — veg soft-argmax blew up ( $RMSE_{vH}$  7.0 vs 2.9), no building gain. Shared grid stays best.



## Design 4 — Two-stage training

---

## Two-stage: train, then purify (1 height + 2 seg)



**Stage 1** — coupled, 80 ep  
AdamW lr  $2 \times 10^{-4}$ ; train everything  
jointly  $\Rightarrow$  a *shared init*.

Seg (crisp, IoU) and height (smooth, RMSE) fight over the shared backbone, so each purify freezes one task's gradient and lets the other own it (a *mirror*).

**Stage 2** — purify off stage 1, 20 ep, lr  $1.5 \times 10^{-4}$

- height-purify (presence-trunk-grad-scale 0)  $\rightarrow$  **ch 3**
- seg-purify (height-trunk-grad-scale 0), then +20 ep cIDice on top  $\rightarrow$  **ch 0-2**

**Fits:**  $\sim 50$  ep already reproduces the 80-ep result within noise (Stage 1 keeps its *best-val* checkpoint).

**Loss**

---


## Loss — every term tracks the metric

ImprovedCompositeLoss (presence-centered):

| Term                                   | Target               | Metric          |
|----------------------------------------|----------------------|-----------------|
| Presence BCE (ring-weighted)           | presence logits      | IoU             |
| Presence Tversky ( $\alpha$ per class) | presence prob        | IoU prec/recall |
| Height L1 + bin-CE                     | soft-bin specialists | RMSE            |
| Building-boundary BCE+Dice             | aux edge head        | IoU edges       |
| Water empty top- $k$                   | water logits         | water IoU (FP)  |

Weights: bld height boost  $\times 5$ , veg  $\times 1.5$ ; bin-CE 0.5; bg weight 0.05; ring  $\alpha=2$  (k5); water top- $k$  512 @ 0.03; boundary 0.25.

## Height loss: L1 + bin-CE




L1 weight  $1 + 5m_{\text{bld}} + 1.5m_{\text{veg}}$

Heights normalised /30 m. Two terms:

- **L1**  $|\hat{h} - h|$  on the soft-argmax height: blended output (weights, left) + per-class aux on own-class pixels.
- **bin-CE** (0.5): Gaussian soft-target over bins ( $\sigma=1.5$ ) vs. predicted bin distribution → *commit* to the right bin, not collapse to a safe mean. (the key soft-bin term)
- **Base** specialist: bin-CE over *all* valid pixels + blended L1 → bg/water height ( $\approx 0$ ) is still constrained, just not metric-scored.

Masks use `gt_class>thr` — the same per-class mask the RMSE metric scores on.

# Ring-weighted presence BCE



dashed = dilate / erode (k5); solid = mask  
edge

- **ring** =  $\text{dilate}(m) - \text{erode}(m)$ : a  $\pm 2\text{px}$  band on the building edge.
- weight  $w = 1 + 2 \cdot \text{ring} \Rightarrow$  ring pixels count  $3\times$ , all else  $1\times$  (only the building channel).
- **Fits**: 88% of lost building IoU sits in this band; equal weighting lets the easy interior/bg dominate.
- Reweights the *main* presence logit (what IoU reads), not a separate boundary head.

# clDice — a topology loss for merged buildings

## The failure it targets

- Buildings are 1–2 px wide; the model *merges* adjacent ones (predicts 67 components/tile vs. GT's 104).
- **51% of building FP** is inter-building “bridges”; oracle un-merge  $\rightarrow$  IoU 0.504  $\rightarrow$  0.587.

## clDice (cl-dice-weight)

- `soft_skeleton`: a differentiable skeleton via iterated min/max-pool (no parameters), then centerline Dice on the building presence map.
- rewards *connectivity*, penalises merges — orthogonal to the area (Tversky) losses.
- applied only as a late 20-ep purify (not in stage-1).

**Not in stage-1:** baking clDice into stage-1 pulled the backbone toward topology and *degraded height* ( $\text{RMSE}_{bH} + 0.006$ ); the late-purify placement keeps height untouched.

## Inference & Results

---



- **5-fold, leave-region-out CV** — so local val predicts the public board ( $LB \approx$  worst fold).
- **Mean ensemble** across folds; optional 0.3 water-expert blend.
- **Per-class threshold sweep** on OOF, then binarize (bld 0.775 / veg 0.675 / water 0.725).
- **Fits:** IoU is judged at a hard cut, different per class  $\rightarrow$  tune each on held-out data.

Ruled out (tried, no gain): TTA, D4 aug, bigger ensembles, region-conditioning.

## Final model: xf095\_2stage\_softbin\_covgt10

Two-stage dual-trunk + shared soft-bin. 5-fold (seed 0), purify checkpoints:

| fold        | Score        | iou_bld | iou_wat | RMSE <sub>bH</sub> | RMSE <sub>vH</sub> |
|-------------|--------------|---------|---------|--------------------|--------------------|
| f0          | 0.509        | 0.511   | 0.455   | 1.60               | 2.92               |
| f1          | 0.541        | 0.512   | 0.466   | 1.40               | 2.70               |
| f2          | 0.521        | 0.497   | 0.448   | 1.48               | 2.78               |
| f3          | 0.530        | 0.500   | 0.465   | 1.43               | 2.74               |
| f4          | 0.525        | 0.522   | 0.468   | 1.48               | 2.93               |
| <b>mean</b> | <b>0.525</b> |         |         |                    |                    |

Worst fold f0  $0.509 \approx$  public LB. Building IoU still boundary-limited (88% of lost IoU in a  $\pm 2\text{px}$  band); that is the open lever.

# Final Submission

---

## From one model to the final ensemble

Everything above is *one* model. The submitted system ensembles **5 architectures**  $\times$  **5 folds**, each run through **3 training stages**:

| Member                      | pixel_backbone | seeds   |
|-----------------------------|----------------|---------|
| U-Net++ (nested decoder)    | unetpp         | 0, 1, 2 |
| UNet 3+ (full-scale skip)   | unet3plus      | 0       |
| TransUNet (attn bottleneck) | unetpp_trans   | 0       |

- **Fits:** identical inputs  $\Rightarrow$  *architecture* diversity decorrelates errors better than more seeds of one net.
- 5-fold leave-region-out; the ensemble is a plain mean of member-fold probability maps, then thresholded.

U-Net++ was the only backbone that beat plain LightUNet on the official IoU metric; heavier/attention encoders all lost.

## Seg from two decorrelated sources (50-member)

Final channels:

- **seg ch 0–2** = mean of **50** member-folds = 25 clDice  $\cup$  25 seg-purify (the two share a stage-1 but diverge in the seg fine-tune  $\Rightarrow$  decorrelated).
- **height ch 3** = mean of **25** height-purify member-folds *only* (adding stage-1 height *hurt* — it is strictly worse).

**Fits:** pure variance reduction — no parameter is tuned on the public board, so it transfers to the private split. On public it lifted all three IoUs (bld +0.0013, veg +0.0011, water +0.0015).

## Height calibration: explaining the local→public gap

Local (OOF, tile-RMSE) vs. public, *no shift*:

|                    | OOF tile | OOF pixel | public |
|--------------------|----------|-----------|--------|
| $\text{RMSE}_{bH}$ | 1.48     | 2.32      | 1.81   |
| $\text{RMSE}_{vH}$ | 2.74     | 2.97      | 3.08   |

- **Convention is tile-based:** the tile gap is a *consistent* +0.33 m on *both* heights (region shift); the pixel gap is inconsistent (building would improve — impossible).
- **Mechanism:** soft-argmax *range compression* (tall under-predicted)  $\times$  a taller test region  $\Rightarrow$  systematic under-prediction.

## Height correction: per-class calibration (applied in final)

### Building: de-compression (used)

- $h \rightarrow 1.05h + 0.116$  (slope  $> 1$ , un-squeeze the compressed tall tail).
- motivated by the strong compression curve ( $-4.5$  m @ 15 m).
- region-invariant  $\Rightarrow$  private-safe.  $\text{RMSE}_{bH}$   $1.807 \rightarrow 1.782$ .

### Vegetation: additive shift (used)

- $h \rightarrow h + 0.12$  on predicted-veg pixels (pure mean offset).
- corrects the small public veg bias.  $\text{RMSE}_{vH}$   $3.079 \rightarrow 3.072$ .

Both corrections are region-invariant (a mean offset, plus a mild slope for buildings), not public-fitted; the residual  $+0.33$  m is irreducible region-shift scatter.

## Final result & trajectory

| Submission               | Score         | bld          | veg          | wat          | RMSE <sub>bH</sub> | RMSE <sub>vH</sub> |
|--------------------------|---------------|--------------|--------------|--------------|--------------------|--------------------|
| 3-seed U-Net++           | 0.5018        | 0.498        | 0.820        | 0.521        | 1.807              | 3.079              |
| 5-member arch ensemble   | 0.5026        | 0.498        | 0.819        | 0.526        | 1.804              | 3.079              |
| + seg-purify + bld shift | 0.5061        | 0.502        | 0.821        | 0.525        | 1.778              | 3.079              |
| + cIDice + veg shift     | 0.5063        | 0.502        | 0.820        | 0.526        | 1.778              | 3.072              |
| <b>+ 50-seg (final)</b>  | <b>0.5067</b> | <b>0.503</b> | <b>0.821</b> | <b>0.527</b> | 1.782              | 3.072              |

- **+0.0049** over the session; the gains are robust (variance reduction + a region-invariant mean shift), not public-fitted.
- Private = best-public auto-selected  $\Rightarrow$  near the end only *non-public-tuned* levers were used.



# Ablation log — what moved the needle, what didn't

## Kept (each +):

- cross-architecture ensemble
- seg-purify (dedicated seg backbone)
- cIDice topology loss
- per-class height calibration (bld de-comp  $1.05h+0.116$ , veg  $+0.12$ )
- 50-member dual-source seg
- delmask (drop human-deleted footprints)

## Refuted (tried, no gain):

- ViT / SegFormer / HRNet / UPerNet / UNet3+ (solo)
- wavelet down-sampling; deep supervision
- cIDice-from-stage-1
- per-code thresholds; morphology; valley-cut
- distance-watershed (1–2px  $\Rightarrow$  degenerate)
- AdaBN; region-conditioning; TTA; D4 aug

Building IoU is at the *embedding-resolution* ceiling (1–2px buildings); the gap to the top teams is a structural model-quality difference (better building-seg / veg-height models on identical embeddings).

## Summary

---

## Summary — each design answers “what is the input?”

| Design                      | Driven by...                                 |
|-----------------------------|----------------------------------------------|
| Dual U-Net++ + gate         | dense $256^2$ , differing reliability        |
| Cross-source token FiLM     | coarse $16^2$ context (modulate, not pixels) |
| Selective calib + zero-init | wild scales (Thor); optional correction      |
| Split-trunk head            | IoU-crisp vs. RMSE-smooth conflict           |
| Presence-gated height       | per-class RMSE mask                          |
| Two-stage purify            | conflict reaches the backbone                |
| Per-class threshold sweep   | IoU judged at a hard per-class cut           |